

An introduction to SDRAM and memory controller

芯片平台-ASIC

邢志国

2019-01



目录

CONTENTS

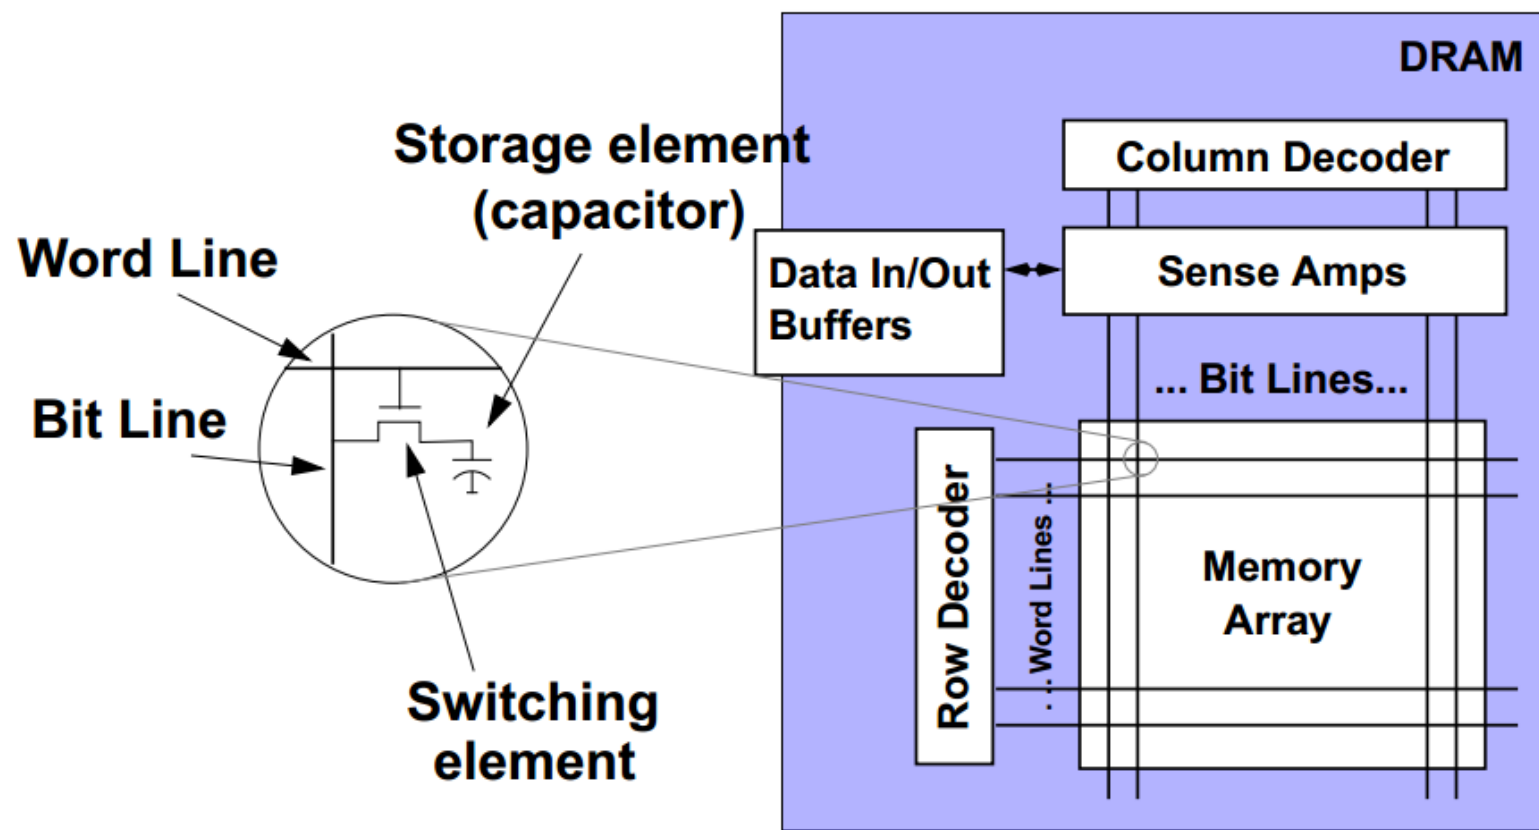
1. DDR SDRAM Brief Introduction
2. DDR SDRAM Efficiency
3. LPDDR3 VS LPDDR4
4. DDR Memory Controller

Part 1

DDR SDRAM Brief Introduction

Dynamic RAM (DRAM)

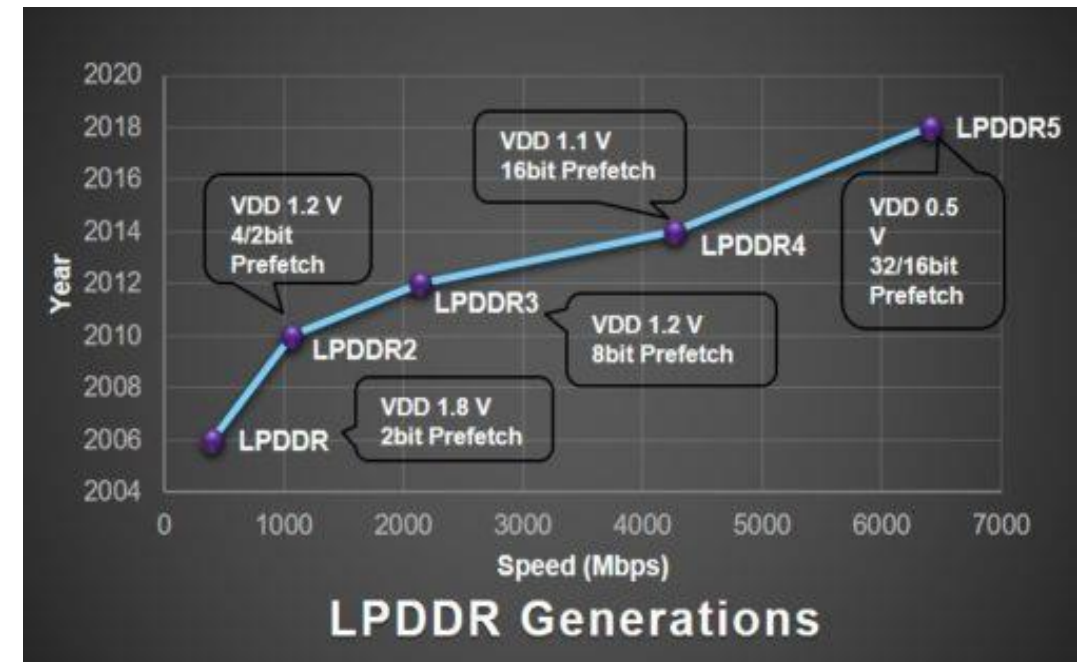

DRAM was patented in 1968 by **Robert Dennard** at IBM



Evolution of the DRAM Design


SDRAM (single Data rate): A clock signal was added making the design synchronous

DDR (Double Data Rate): The data bus transfers data on both rising and falling edge of the clock



DDR5:

- DDR4 successor JEDEC standard for Server/Enterprise markets
- Targeting data rates up to 6400Mbps
- DDR5 samples projected for H1/2019




LPDDR5:

- LPDDR4/4X successor JEDEC standard for Mobile/ Automotive markets
- Targeting data rates up to 6400Mbps
- LPDDR5 samples projected for mid 2019

DDR(Double Data Rate)

- Double-data-rate architecture; two data transfers per clock cycle
- **Differential clock inputs**
- Bidirectional, data strobe (DQS) is transmitted/received with data, to be used in capturing data at the receiver
- DQS is **edge-aligned** with data **for READs**; **center-aligned** with data **for WRITEs**
- Read and write accesses to the DDR SDRAM are **burst oriented**.



Prefetch :

Internal data bus width = Prefetch value * external bus width


DDR : 2n

DDR2/LPDDR2: 4n

DDR3/LPDDR3: 8n

DDR4: 8n

LPDDR4/LPDDR4x: 16n



SDRAM Architecture Basic Conception

Row (page)

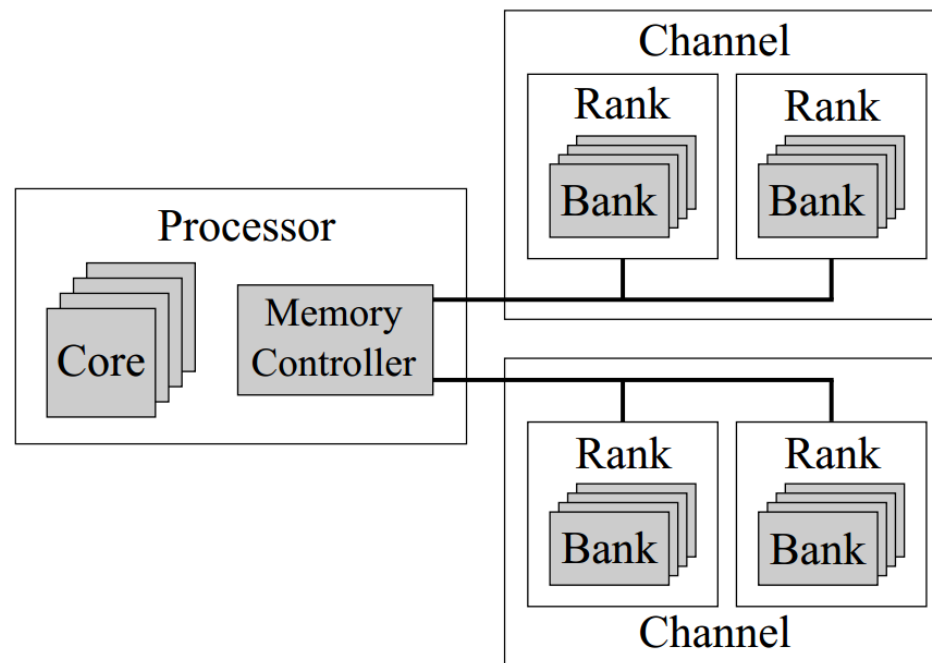

Column

Bank : a bank consists of multiple rows and columns

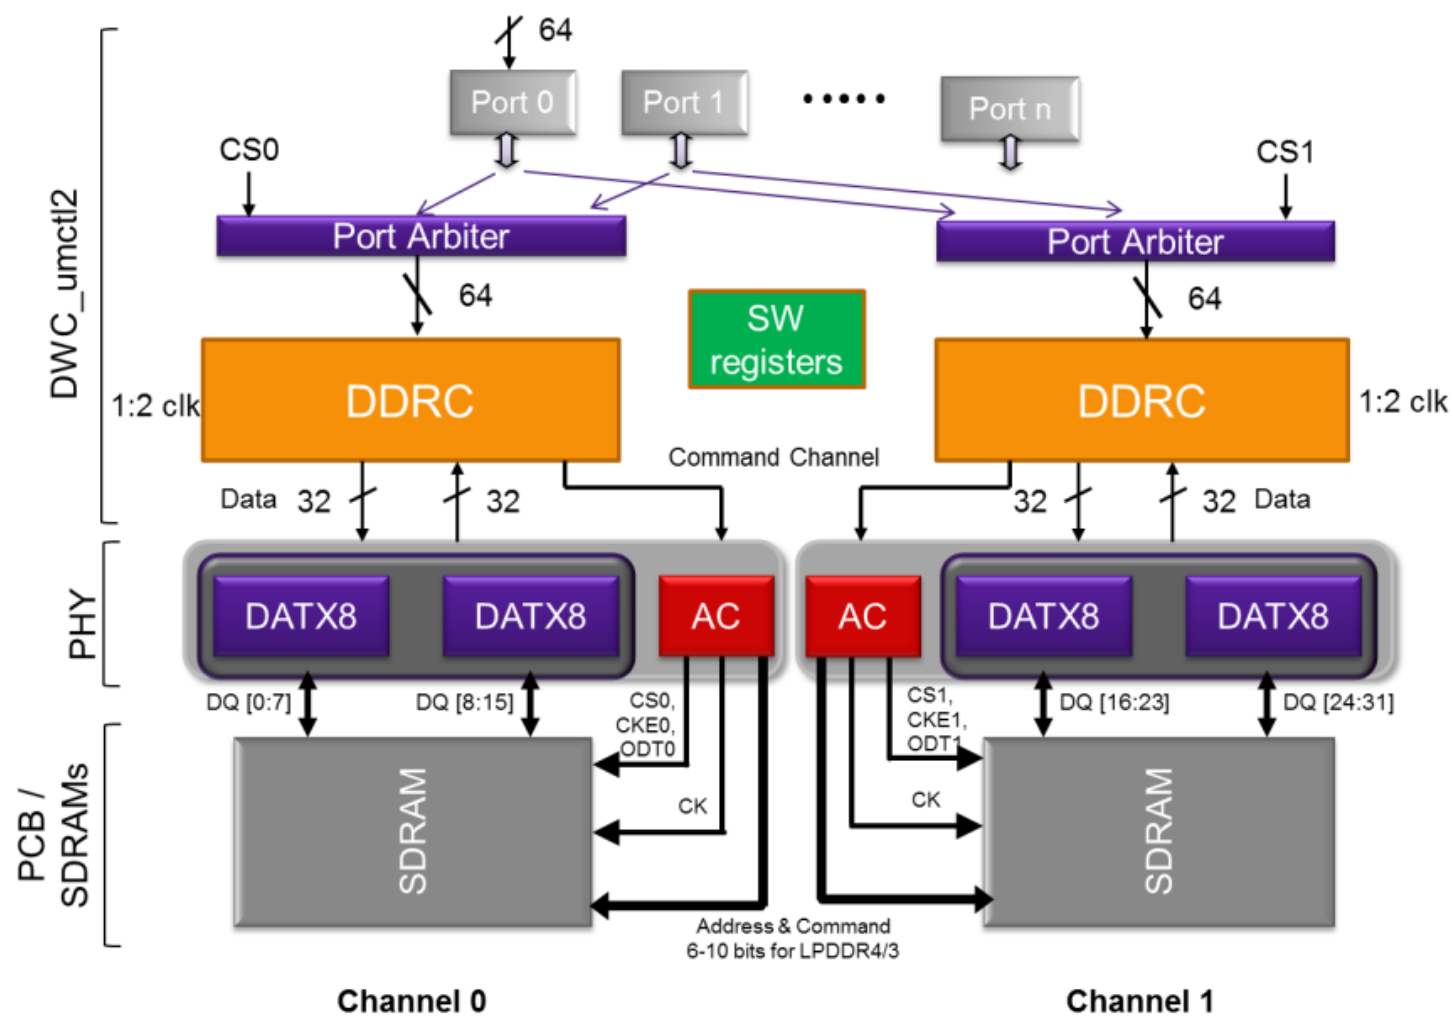
Rank: a set of DRAM chips connected to the same chip select

Channel: Bus --> Controller --> PHY --> SDRAM

Burst Length (BL): determined by prefetch architecture



Example : Dual-Channel



SDRAM Command Summary

Write Command (WR) : Initiate a write burst to an active row

Read Command (RD) : Initiate a read burst to an active row

Mode Register Read (MRR):


Mode Register Write (MRW):

Activate (ACT) : Activate a row in a particular bank

Precharge (PRE): Close a row in a particular bank independently or rows of all banks simultaneously

Refresh (REF) : Start a refresh operation

NOP/DES : No Operation / Ignores all inputs



Refresh Operation

Refresh :

Refresh is the periodic command required by DRAM to maintain the charge on memory cells in the DRAM device.

The exact refresh interval time depends on the DRAM type (e.g., DDR or LPDDR) and the operating temperature.


Auto-refresh:

Auto refresh is generally a term associated with automatic generation of refresh commands with the memory controller;

Self-refresh:

The DRAM supports a mode where refresh is maintained internally by the memory device. Self refresh is a low power mode and generally entered when memory is IDLE. No read or write transactions can occur when memory is in self refresh.

DDR SDRAM initialization (Example : LPDDR4)



Memory Training

Command /Address Training (CA Training) : Centers the DRAM clock in the center of the DDR CA data eye.
LPDDR3/LPDDR4 Only

Write leveling : align DQS rising edge with the MEMCLK rising edge

Read Gate Training (aka RxEn training) : determine the round-trip latency

Read DQ deskew training : compensates for delay differences among the DQ lanes, primarily caused by board routing and DRAM DQ output skew

Data Eye Training (Read Leveling): center the read DQS timing in the read DQ passing region.

VREF Training :Defines the reference voltage range. VREF training is only applicable for DDR4/LPDDR4 memory systems

Part 2

DDR SDRAM Efficiency

Memory efficiency

Memory efficiency is the fraction between the amount of clock cycles when data is transferred and the total amount of clock cycles

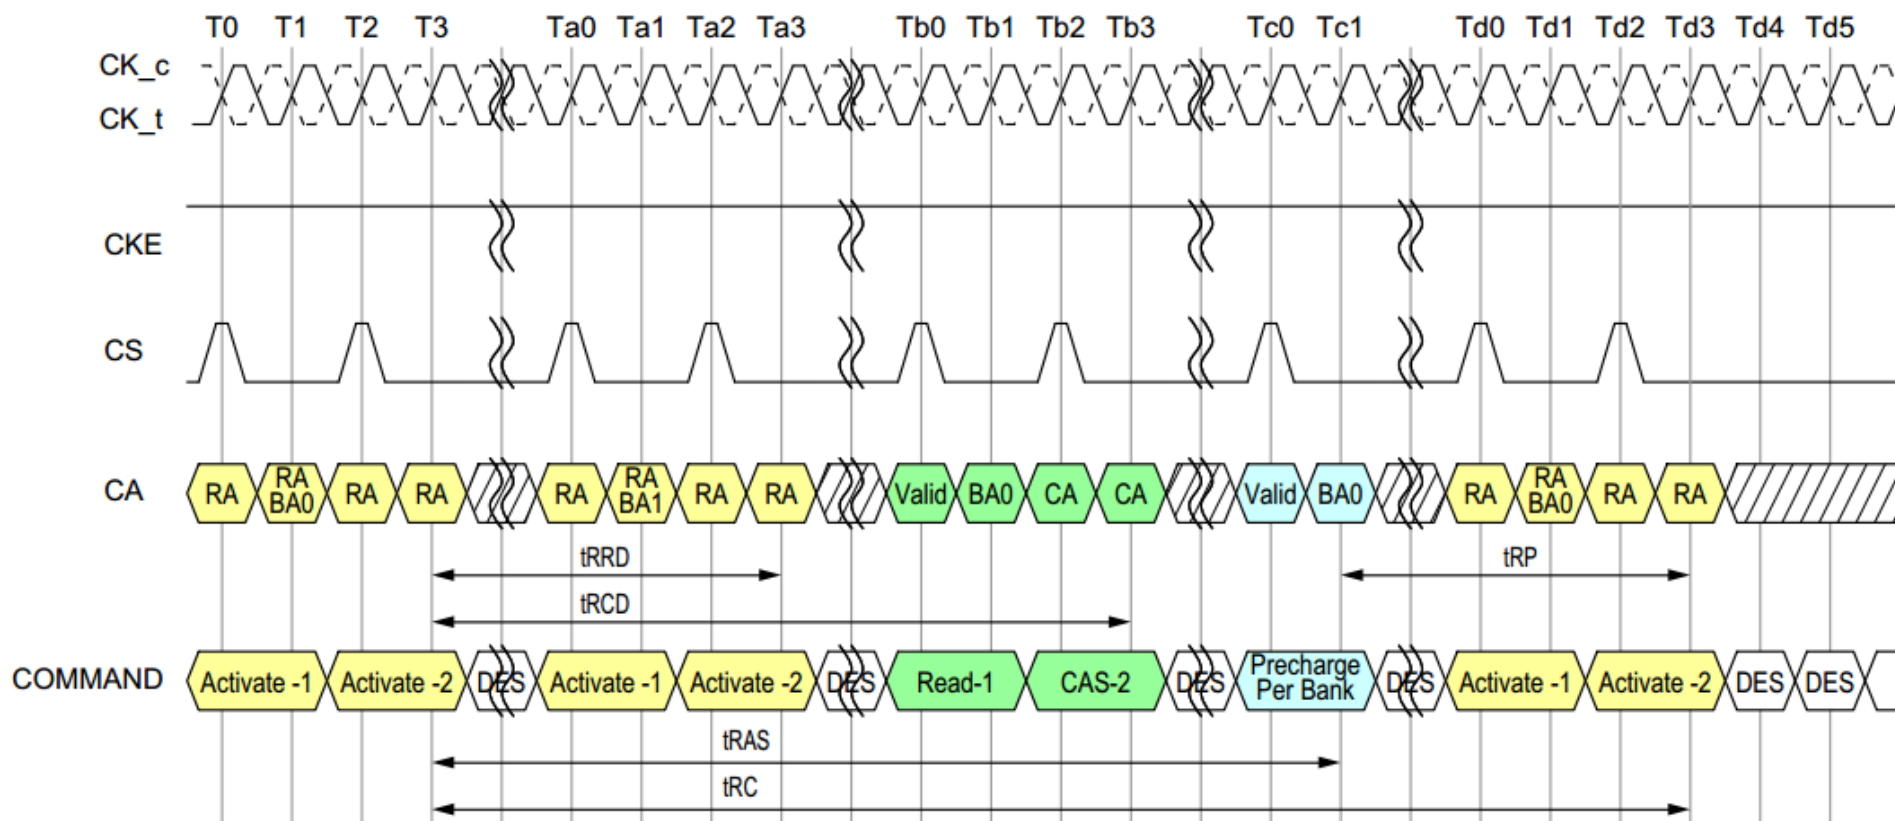
- Defines the exchange rate between peak bandwidth and actual bandwidth
- $\text{peak bandwidth} = \text{memory clock freq} * 2 * (\text{memory data bus width}/8)$

Four categories of memory efficiency for SDRAM:

- ◆ Bank conflict efficiency
- ◆ Refresh efficiency
- ◆ Data efficiency
- ◆ Read/write efficiency

Bank conflict efficiency

Bank conflicts : two consecutive bursts to different rows in the same bank.



Timing Constraints (LPDDR4 3200Mbps)

Parameter	Definition	Cycles(tCK)
tRRD	Active bank-A to active bank-B	16
tRCD	RAS-to-CAS delay	29
tRP	Row precharge time	29(per bank)
tRAS	Row active time	68
tRC	ACTIVATE-to-ACTIVATE command period (same bank)	97(per bank)

Refresh efficiency

Refresh requirement:

- A bank must be in the idle state before it is refreshed
- No data can be transferred when the memory is being refreshed.

Degrade performance:

Any accesses to a DRAM bank that is refreshing must wait for the refresh latency $tRFC$,

All bank refresh : operates at the rank level, refreshes a number of rows in all banks of a rank concurrently.

Per bank refresh : enhanced mode for LPDDR, This enables a bank to be accessed while another in the same rank is being refreshed,

Data efficiency

A memory burst cannot access segments of the minimum burst length

Minimum access granularity

- Burst length 16, words is 64 B with 32-bit memory

If data is poorly aligned an extra segment has to be transferred

Cycles are lost when transferring unrequested data



Read / write efficiency

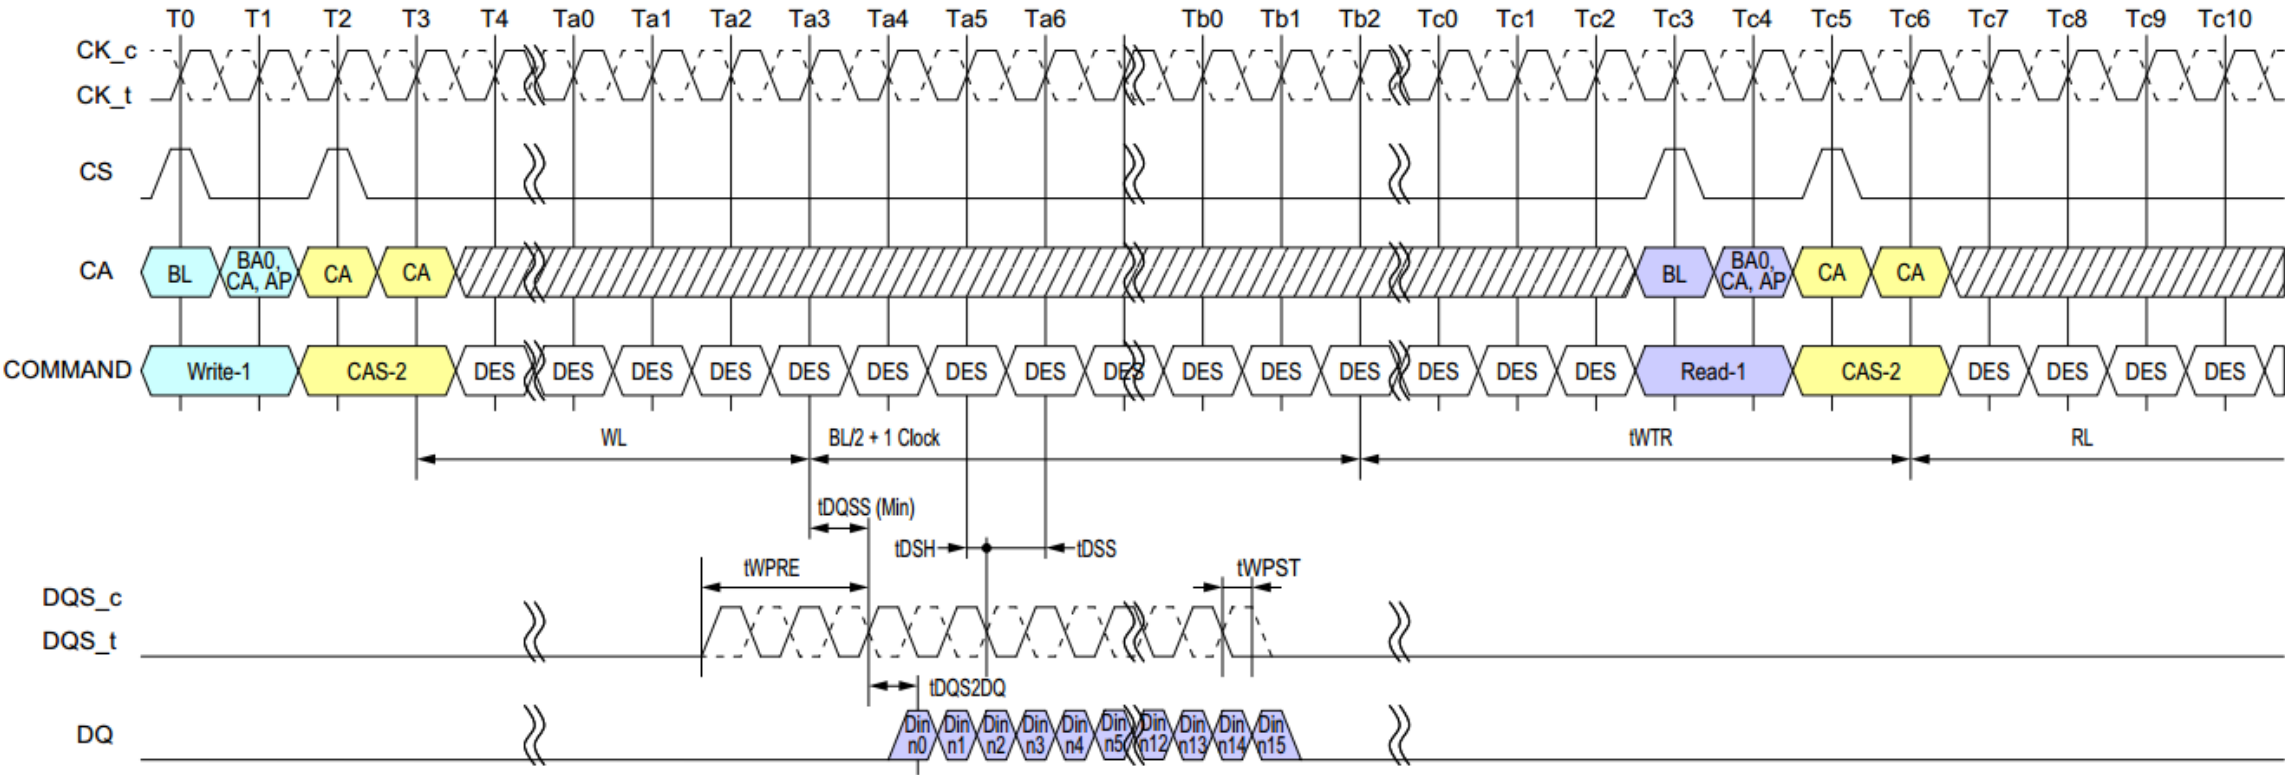
The data bus of an SDRAM is bi-directional

- ◆ Cycles are lost when switching direction of the data bus
- ◆ Extra NOPs must be inserted between read and write commands

Read/write efficiency depends on traffic

- ◆ Determined by frequency of read/write switches
- ◆ Switching too often has a significant impact on memory efficiency

Read / write efficiency




Burst Write Followed by Burst Read

23

Part 3

LPDDR3 VS LPDDR4

LPDDR3 VS LPDDR4





LPDDR3 VS LPDDR4

■ Evolutionary DRAM technology enables 3.2Gbps and faster

Items		LPDDR3	LPDDR4	Comments
Speed	CLK	400-800MHz (~1066MHz w/ LP3E)	800-1600MHz	2X, Pursues higher speed
	CMD/ADDR	DDR	SDR	-
	DQ	DDR	DDR	
	Band Width	12.8GB/s+ (2ch)	25.6GB/s+ (2ch)	2X
Voltage	VDD2/VDDQ/VDD1	1.2/1.2/1.8	1.1/1.1/1.8	Total Pd 10%↓
Architecture	[# o Ch & DQs]/ Die	x32	2x16	IDD4 20% ↓
	# of Bank/channel	8	8	-
	Page Size	4K	2K	IDD0 10%↓
	BL	8	16	32B/ch
Interface	I/O interface	HSUL	LVSTL	40% I/O power reduction (vs. POD)
	DQ ODT	No term (VDDQterm option)	VSSQ Term	
	CA ODT	No term	VSS term	Optional
	Vref	External	Internal	


LPDDR3 VS LPDDR4

Low Power 32-bit Architecture




~20% IDD4 reduction

Access Granularity & Page-size/Ch



LPDDR3 : Activated Page 4KB(1Kx32)



LPDDR4 : Activated Page 2KB(1Kx16)

~10% IDD0 reduction

LPDDR3 VS LPDDR4

Data Bus Inversion (DBI): Data Bus inversion is a power saving DDR4/LPDDR4 specific feature which aims at minimizing:

DDR4: Logic 0

LPDDR4: Logic 1

Simultaneous switching outputs

DBI removes a significant amount of SSO noise from data transactions.

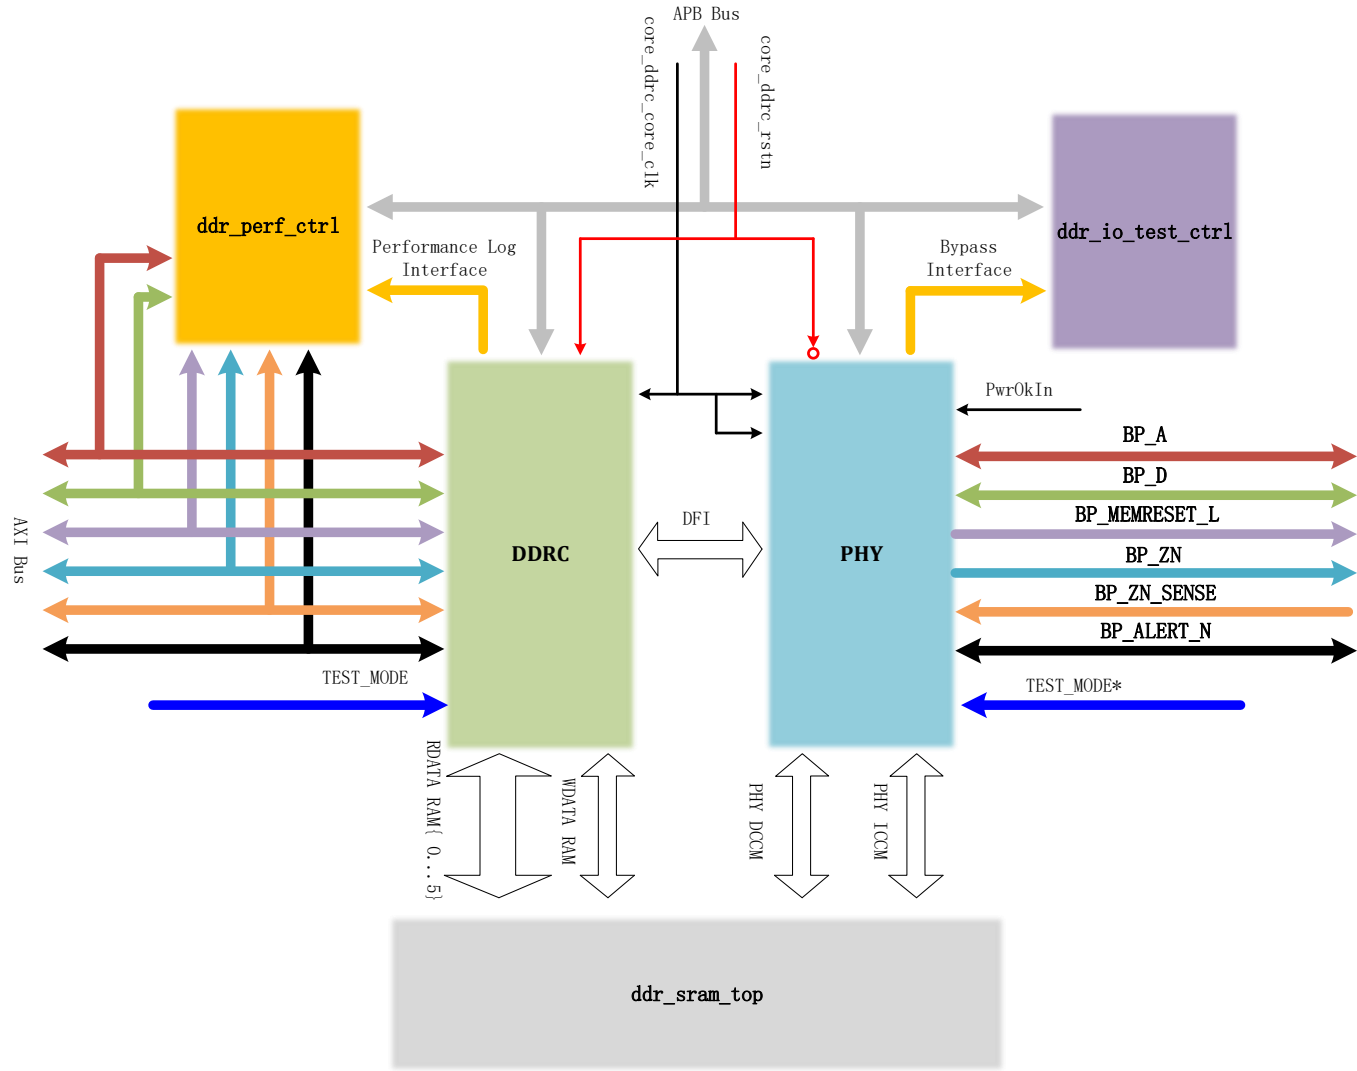
Example : DDR4

DQ[0:7] per UI	DBI	After Converting Data	# of "0" including DBI
0000_0000	0	1111_1111	1
0000_0001	0	1111_1110	2
0000_0011	0	1111_1100	3
0000_0111	0	1111_1000	4
0000_1111	1	0000_1111	4
0001_1111	1	0001_1111	3
0011_1111	1	0011_1111	2
0111_1111	1	0111_1111	1
1111_1111	1	1111_1111	0

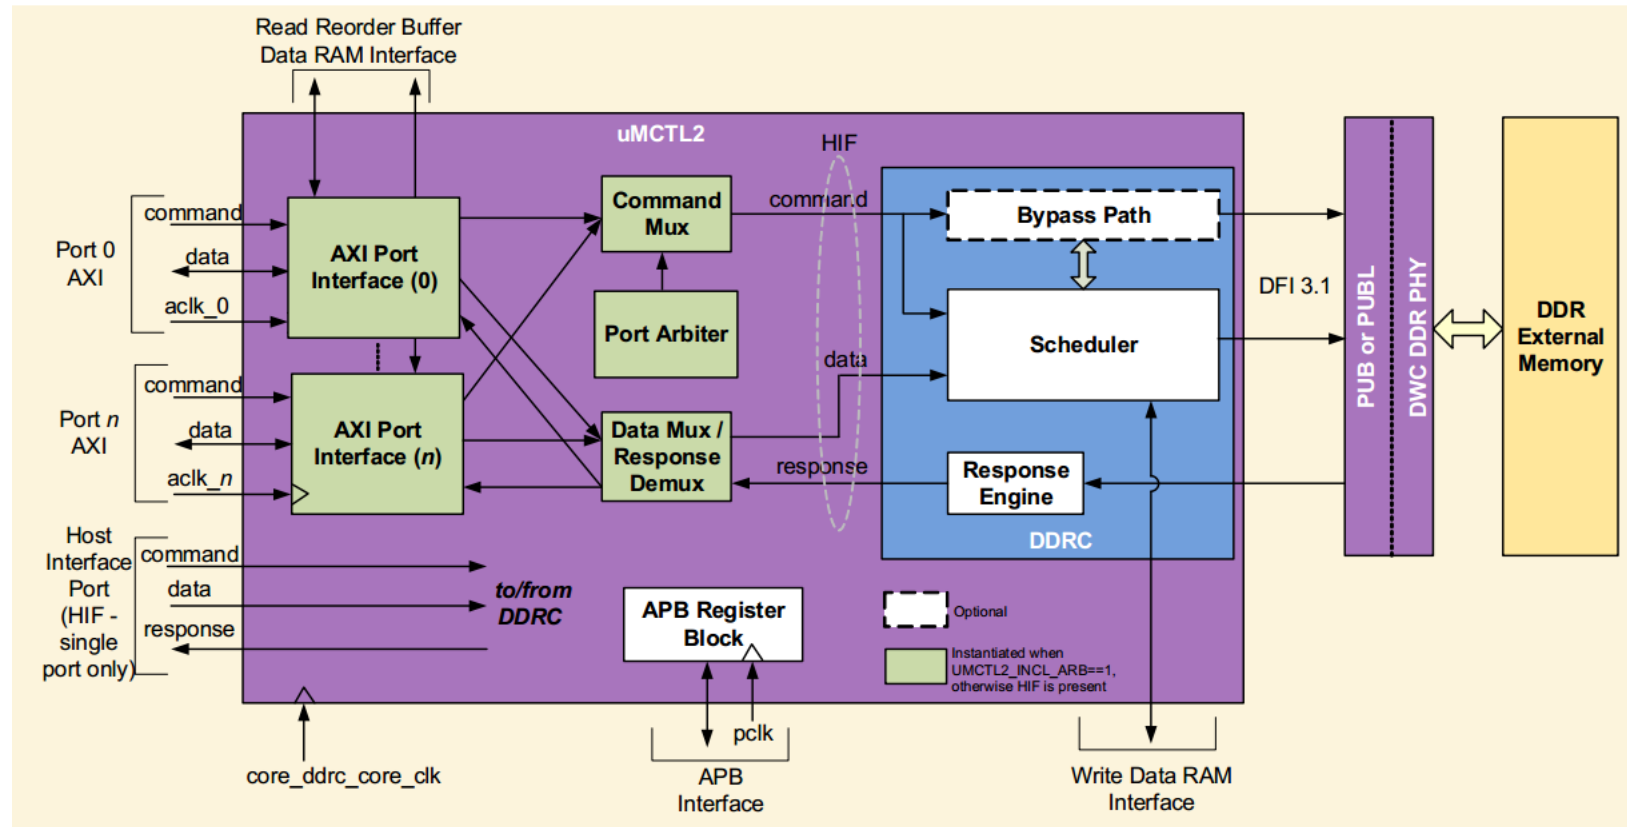
Part 4

DDR Memory Controller

X2 DDR System-Level Block Diagram



DDR Memory Controller Overview



Major Features

- Multiple memory support for DDR4/3/2 and LPDDR4/4X/3/2
- The uMCTL2 is fully compatible with DFI 4.0 Addendum Version 2 to DFI 3.1
- Scalable 1:1/1:2 frequency ratio architecture.
- Up to 16 host ports using AMBA AXI/AHB
- APB interface for the uMCTL2 software accessible registers
- Read and write buffers in fully associative CAMs, configurable in powers of two, from 16 up to 64 reads and 64 writes
- Hardware configurable and software programmable Quality of Service (QoS) support
- Support 1, 2, or 4 memory ranks
- Support self-refresh entry and exit
- Support for dynamically changing clock frequency while in self-refresh
- Flexible address mapper logic to allow application specific mapping of row, column, bank, and rank bits

DDR Memory Controller Performance

- ◆ For maximum SDRAM efficiency, commands are executed out-of-order
- ◆ CAM-based scheduling and reordering can reorder up to 64 read and 64 write commands
- ◆ Read Reorder Buffer ensures coherent and correctly in-order read responses on each AXI ID
- ◆ Low latency – 6 clock cycle minimum, 8 clock cycle nominal latency
- ◆ QoS features improve system-level performance
- ◆ Optional: Strictly in-order controller (uPCLT2) for external schedulers / NoCs

Port Arbiter

QoS overview

Read Classes

Low Priority Read (LPR): (best effort, no guarantees)

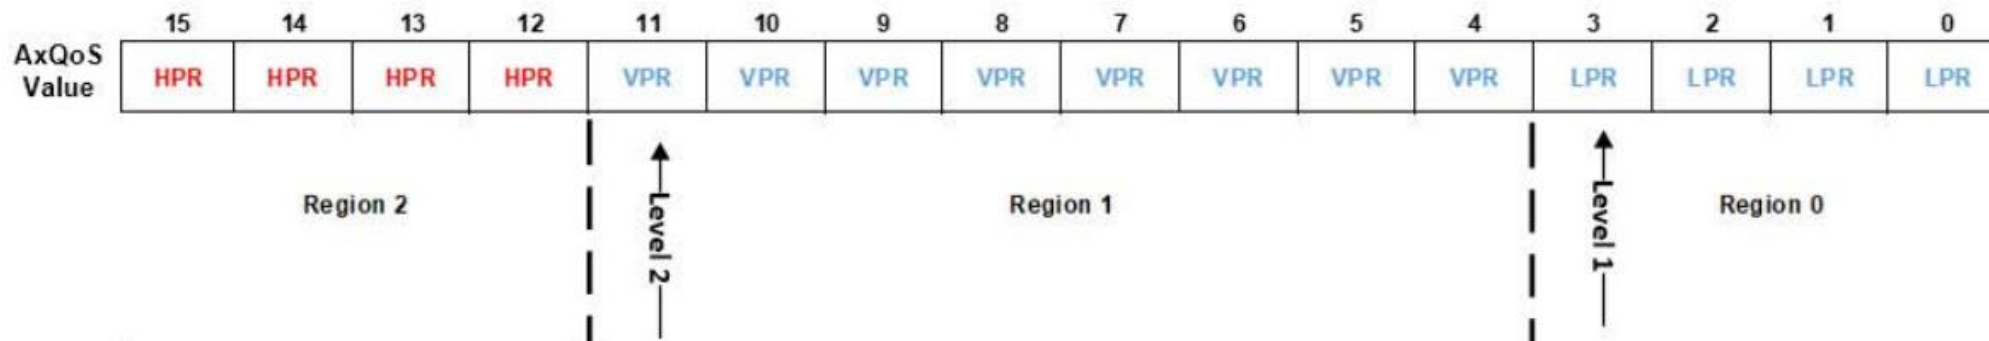
Variable Priority Read (VPR): (maximum latency bound)

High Priority Read (HPR): (minimum latency)

Write Classes

Normal Priority Write (NPW): (best effort, no guarantees)

Variable Priority Write (VPW): (maximum latency bound)



Port Arbiter

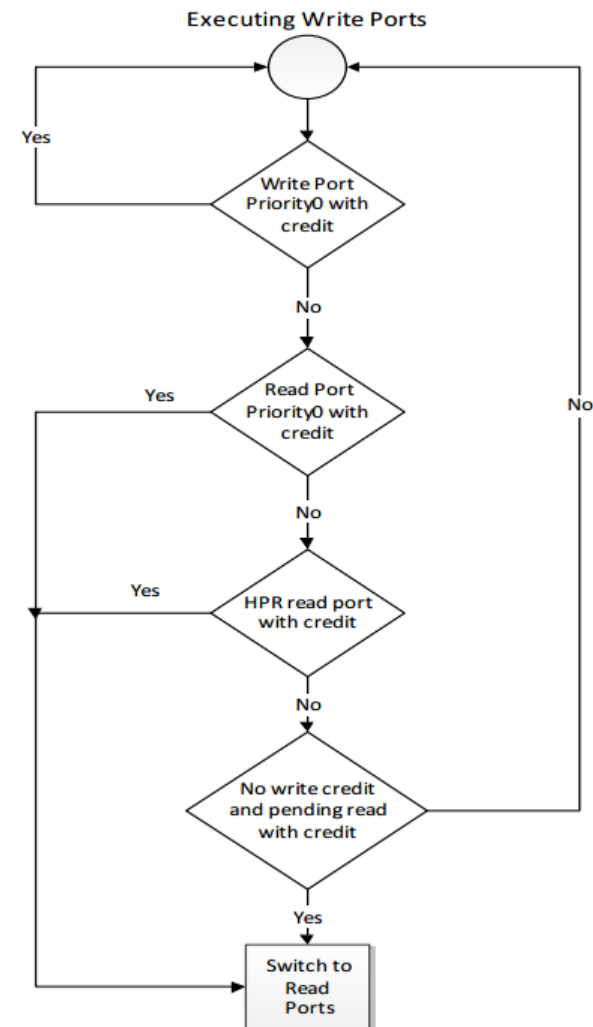
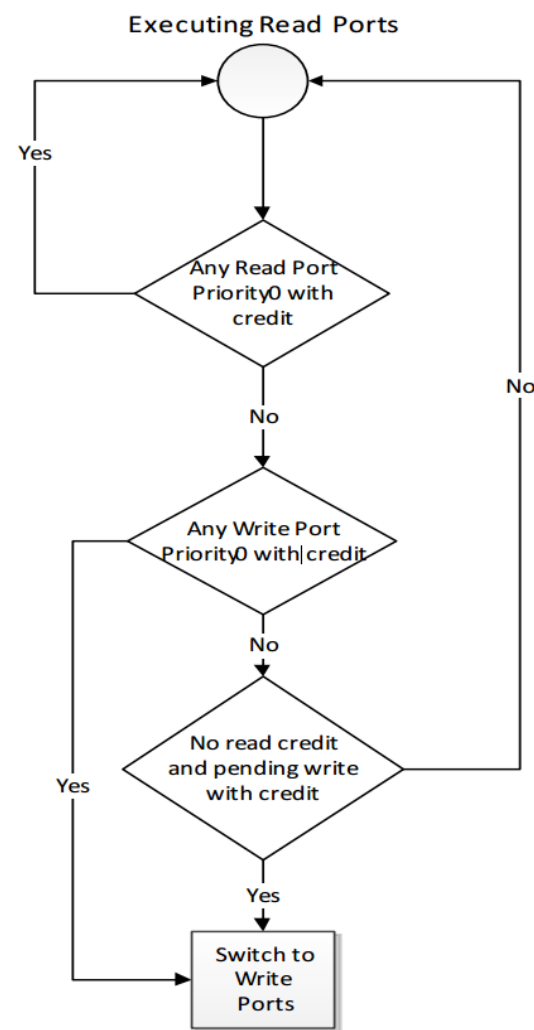
- ◆ Read/Write Arbitration
- ◆ Port Timeout
- ◆ DDRC Read Priorities (HPR/LPR/VPR) and Write Priorities (NPW/VPW) for Ports
- ◆ Port Command Priority
- ◆ Round-Robin Arbitration

Read/Write Arbitration

First level of arbitration , based on:

- Read (LPR-VPR, HPR) and write credits
- Timeout/Expired-VPR/Expired-VPW (Priority0)

If DUAL-HIF enabled , Read/Write Arbitration is not performed by the PA



Port Timeout

- Aging counters (per port, per direction (read/write))
- when a port aging counter becomes 0 and the port becomes the highest priority requester (Priority0)
- The aging feature and the timeout can be enabled by registers

DDRC Read Priorities and Write Priorities

- ◆ initial priority : $HPR > (VPR=LPR)$, $NPW = VPW$
- ◆ VPR (not expired) = LPR
- ◆ VPR (expired) $> HPR$ or writes.
- ◆ VPW (not expired) = NPW
- ◆ VPW (expired) $> NPW$ or HPR .

Port Command Priority

External Port Priorities (UMCTL2_EXT_PORTPRIO = 1)

- The priorities are per-command and can dynamically change for a given port based on AXI AxQoS signals (arqos/awqos).
- There are 16-level priorities set by 4-bit wide arqos/awqos signals where the higher binary value represents a higher priority.

Internal Port Priorities (UMCTL2_EXT_PORTPRIO = 0)

- Another way of setting port priorities is through port aging counters.
- The upper-most 5 bits of the counter determine the port priority out of 32 levels

Round-Robin Arbitration

round-robin : the final arbitration stage.

Cycle	Port Addressed by the Arbitration Counter	Ports Requesting				Command Queue Full?	Arbitration Winner	Value of Counter at Next Cycle
		Port 0	Port 1	Port 2	Port 3			
0	0	Y	Y	Y	Y	Yes	None	0
1	0	Y	Y	Y	Y	No	P0	1
2	1		Y	Y	Y	Yes	None	1
3	1	Y	Y	Y	Y	No	P1	2
4	2	Y		Y	Y	No	P2	3
5	3	Y			Y	No	P3	0
6	0	Y		Y		Yes	None	0
7	0	Y		Y		No	P0	1
8	1			Y		No	P2	2
9	2			Y	Y	No	P2	3
10	3	Y			Y	No	P3	0
11	0			Y		No	P2	1

Transaction Service Control (DDRC Part)

- ◆ Transaction Stores
- ◆ Read/Write Turn-around
- ◆ Address Collision Handling

Transaction Stores

Read Transaction Stores

SCHED.lpr_num_entries register splits the Read CAM in the controller into LPR and HPR sections. The LPR and VPR commands are sent to the LPR section of the CAM, and the HPR commands are sent to the HPR section.

Write Transaction Stores

The NPW and VPW commands are sent to the Write CAM. VPW commands do not have reserved space in the Write CAM.

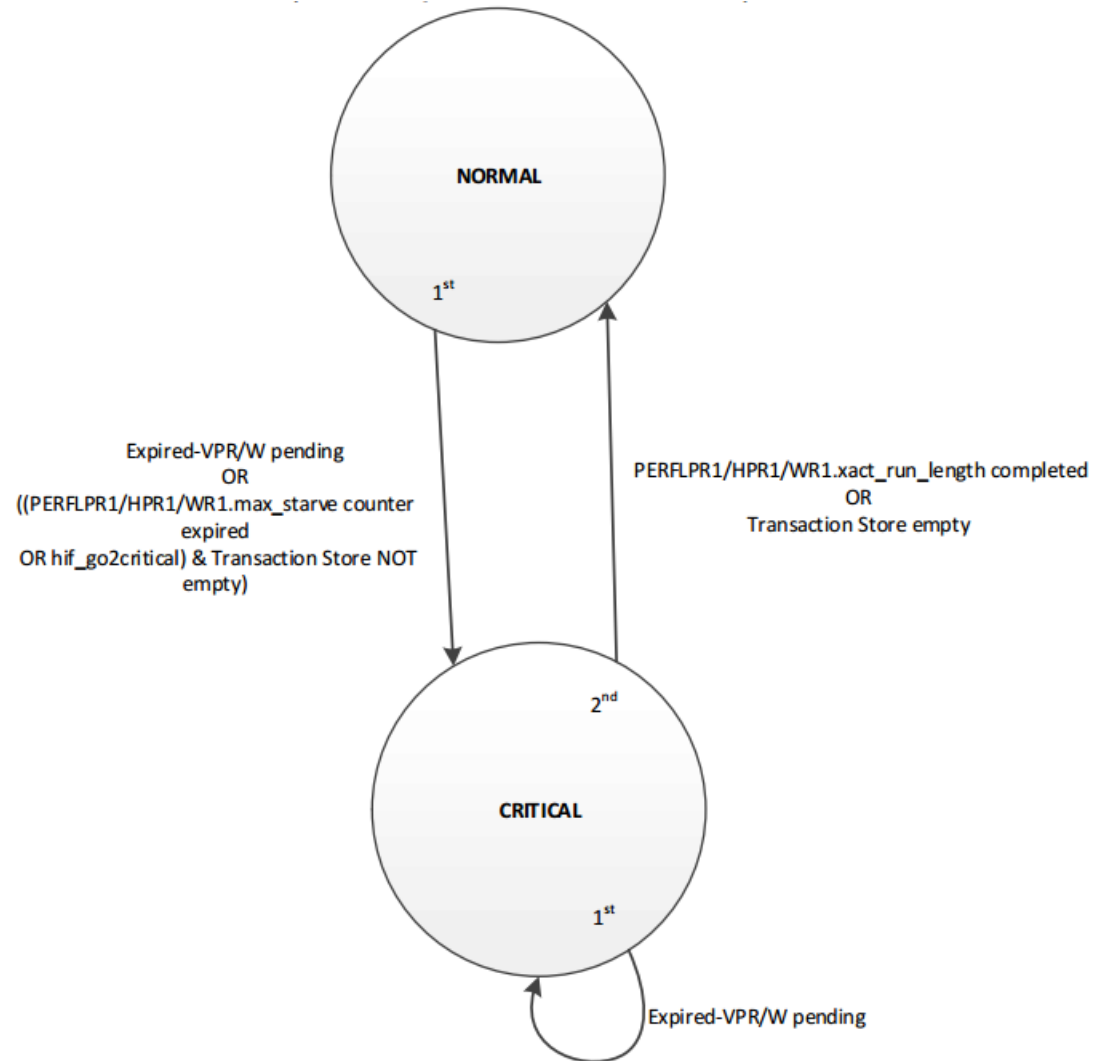
Transaction Store State Transitions

Each class of transaction store (LPR, HPR, or WR) can be in any one of the following two states:

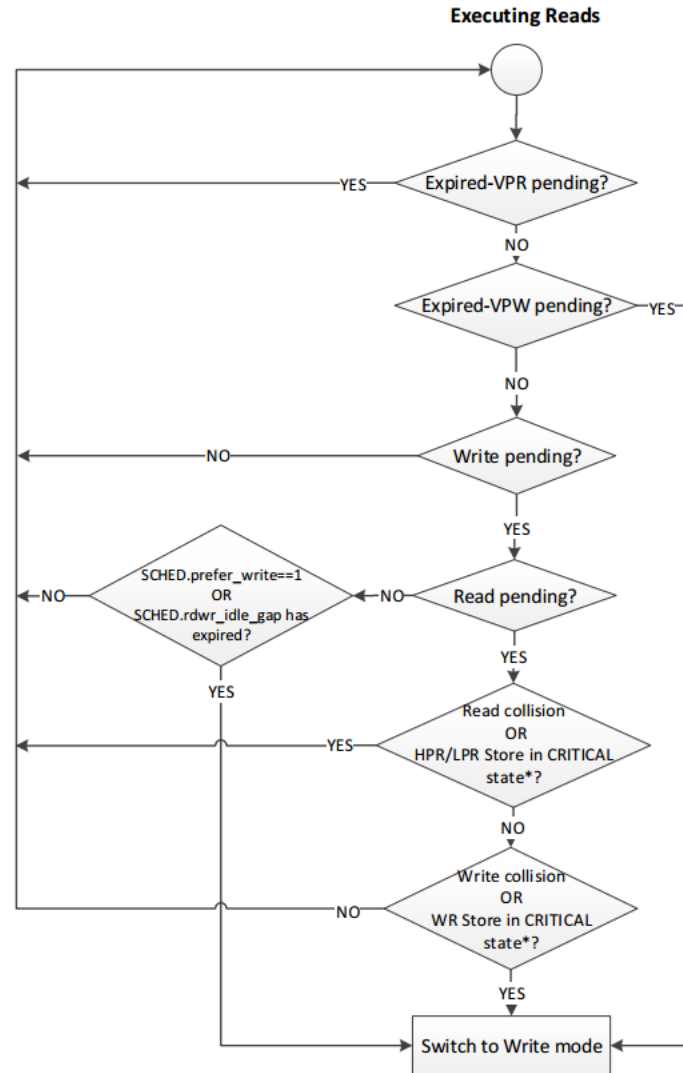
- Normal: It is the state where the transaction store starts.
- Critical: It indicates that the transaction store must be prioritized for service.

The transaction stores move between these two states under the control of *_max_starve and *_xact_run_length registers.

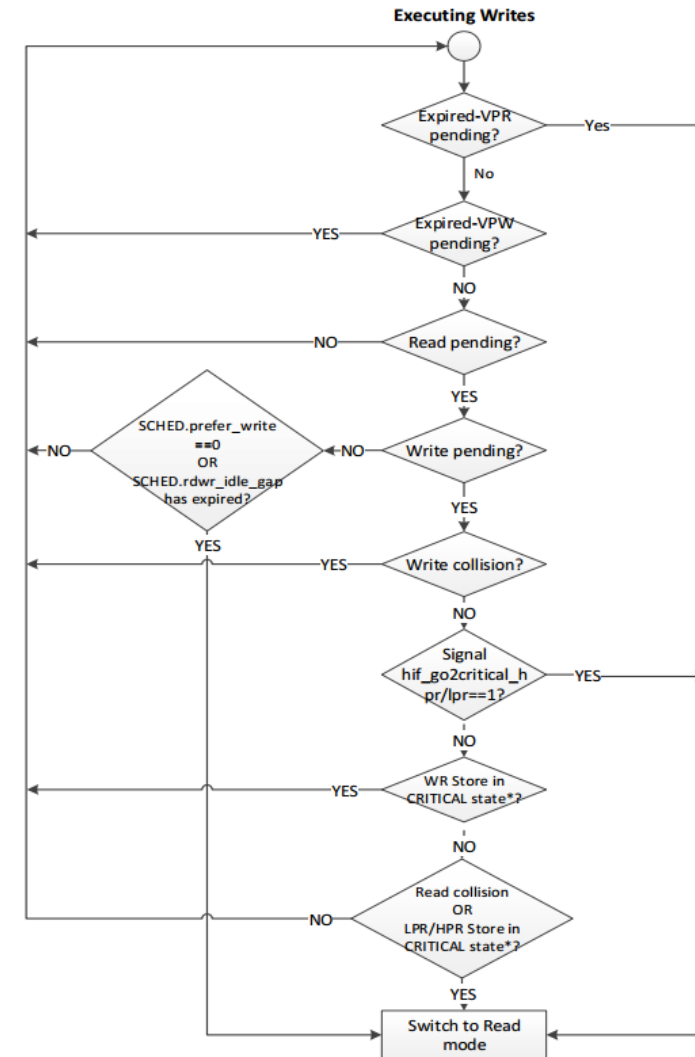
Transaction Store FSM (One FSM per store - WR, LPR, HPR)



Read/Write Turn-around



* CRITICAL state – see Transaction Store FSM diagram



* CRITICAL state – see Transaction Store FSM diagram

Address Collision Handling

New read colliding with queued read: This collision causes no problems. The two reads can end up being executed out-of-order.

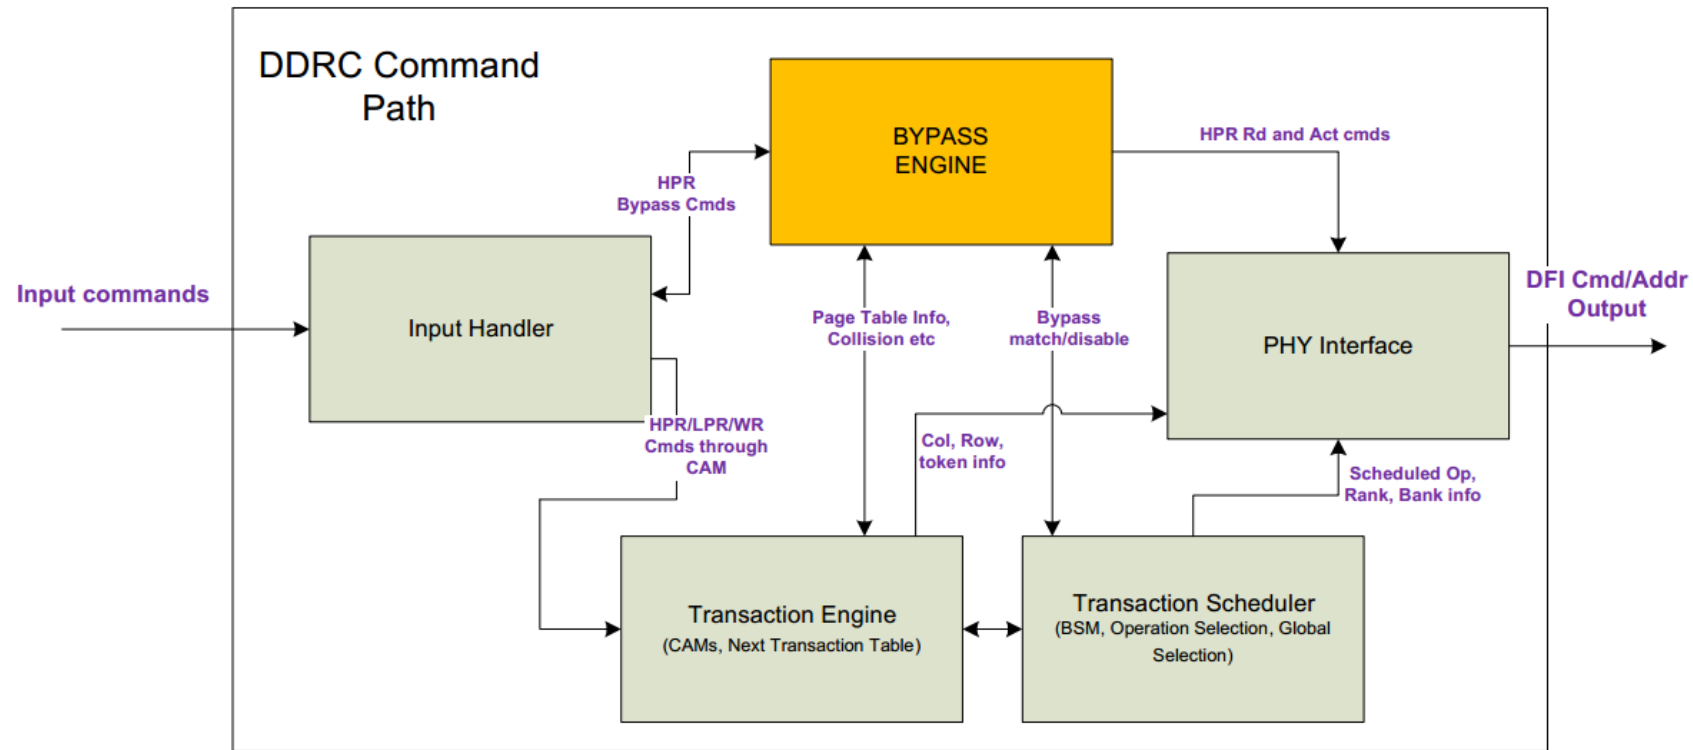
New write colliding with queued write: If write combine is enabled, the DDRC overwrites the data for the old write with that from the new write and only performs one write transaction (write combine).

New read (or write) colliding with queued write (or read) respectively: In this case, the DDRC performs the following sequence:

- a. Holds the new transactions in a temporary buffer.
- b. Applies flow control back to the SoC core to prevent more transactions from arriving.
- c. Flushes the internal queue holding the colliding transaction until that transaction is serviced.
- d. Accepts the new transaction and removes flow control.

Bypass Operation

The bypass engine in the uMCTL2 handles high-priority, low-latency read requests.



Address Mapper

Flexible address mapper logic to allow application specific mapping of row, column, bank, and rank bits

	HIF Address Bits	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
	MEMC_BURST_LENGTH																									
Non-interleaved	8	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	r1	r0	b2	b1	b0	c9	c8	c7	c6	c5	c4	c3	c2	c1	c0
	4	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	r1	r0	b2	b1	b0	c9	c8	c7	c6	c5	c4	c3	c2	c1	c0
Non-interleaved	8	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	r1	r0	b2	b1	b0	c9	c8	c7	c6	c5	c4	c3	c2	c1	c0
Interleaved - example 1	8	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	r1	r0	c9	c8	c7	c6	c5	c4	b2	b1	b0	c3	c2	c1	c0
	4	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	r1	r0	c9	c8	c7	c6	c5	c4	b2	b1	b0	c3	c2	c1	c0
Interleaved - example 1	8	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	r1	r0	c9	c8	c7	c6	c5	c4	b2	b1	b0	c3	c2	c1	c0
Interleaved - example 2	8	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	c6	c7	b2	c8	c9	r0	r1	b1	c3	c5	c4	b0	c2	c1	c0
Interleaved - example 3	4	r11	r10	r9	r8	r7	r6	r5	r4	r3	r2	c9	r0	r1	b0	c7	c8	c5	b1	c6	b2	c4	c3	c2	c1	c0



THANKS!